

Table 1: Configurations of the fine-grained p-chase microbenchmark used to collect the full latency distribution for the NVIDIA TitanX (Pascal).

	L1	N	s	w	r
Configuration 1	✓	256B ($\ll C_{L1}$)	4B	1	1
Configuration 2	✗	256B ($\ll C_{L1}$)	4B	1	1
Configuration 3	✗	6MiB ($\ll C_{L2}$)	128B	0	1